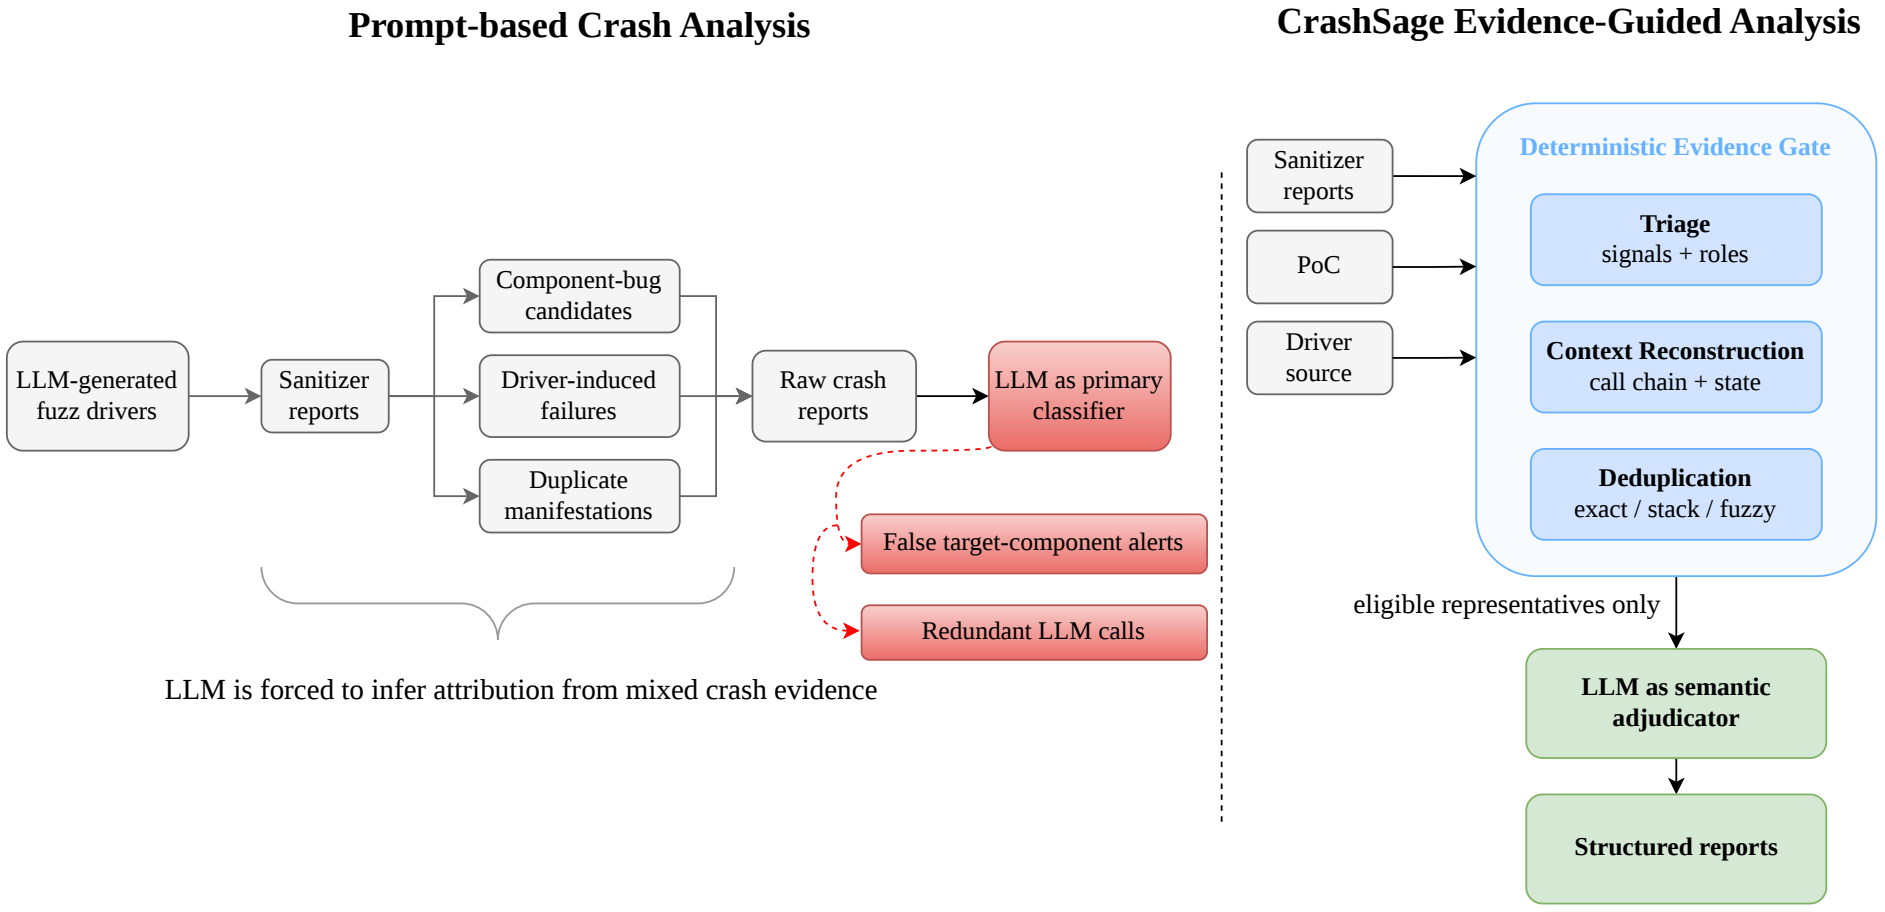
Prompt-based Crash Analysis
LLM-generated fuzz drivers
Sanitizer reports
Component-bug candidates
Driver-induced failures
Duplicate manifestations
Raw crash reports
LLM as primary classifier
False target-component alerts
Redundant LLM calls
LLM is forced to infer attribution from mixed crash evidence
CrashSage Evidence-Guided Analysis
Sanitizer reports
PoC
Driver source
Deterministic Evidence Gate
Triage
signals + roles
Context Reconstruction
call chain + state
Deduplication
exact / stack / fuzzy
eligible representatives only
LLM as semantic adjudicator
Structured reports